

Panda Internet Radio

Goal: Implementing the Panda next_song function based on history of song likes and dislikes.

Side Effect: Learning about docstrings and doctests

Note: This tutorial implements a slightly different spec than defined by the panda.ipynb readme and tested by test.py

Distance Function

A song is represented as a list of genes, it's "genome". Each gene can have value either 0 or 1.

We want a distance function that will give us a sense of how different two songs are.

```
In [1]: def distance(song_0, song_1):
        """A metric indicating 'distance' between two songs

        inputs: Two songs, defined by their genome.
        returns: The "manhattan distance" between two genomes,
                 i.e., the number of genes differing between songs.
        >>> song_0 = []
        >>> song_1 = []
        >>> distance(song_0, song_1)
        0
        >>> song_0 = [0]
        >>> song_1 = [1]
        >>> distance(song_0, song_1)
        1
        >>> distance(song_1, song_0)
        1
        >>> distance([0, 1], [1, 0])
        2
        >>> distance([0], [1, 1])
        Traceback (most recent call last):
          ...
        ValueError: song genomes different length
        """
        #dist = 0
        #for gene in range(len(song_1)):
        #    dist += abs(song_0[gene]-song_1[gene])
        #return dist

        if len(song_0) != len(song_1):
            raise ValueError('song genomes different length')

        return sum([abs(g0-g1) for g0, g1 in zip(song_0, song_1)])
```

```
In [2]: distance([0, 1], [1, 0])
```

```
Out[2]: 2
```

```
In [3]: import doctest
        doctest.run_docstring_examples(distance, globals(), verbose=False)
```

Docstrings and doctests

doctests are test cases embedded within docstrings, that can actually be run and tested automatically:

- Nothing output if all tests succeed
- An error reported if one or more tests fail

These are part of the standard python library, so are always available. Read more about them at

<https://docs.python.org/3/library/doctest.html> (<https://docs.python.org/3/library/doctest.html>)

A useful approach is to include the following at the bottom of your file (e.g., your `lab.py`):

```
In [4]: if __name__ == '__main__':
        # running lab.py invokes the doctests for *all* functions in the file...
        import doctest
        doctest.testmod()
```

Inside jupyter, we'll instead invoke specific doctests directly, e.g.:

```
In [5]: import doctest
        doctest.run_docstring_examples(distance, globals(), verbose=False)
```

Average Distance

Add another data structure -- a "music library" implemented as a dictionary consisting of `song_ids` as keys, and the corresponding genome for the song as the value.

Now we'd like to get a sense of the average distance between one song and a whole list of other songs.

```
In [6]: def average_distance(song_id_list, song_id, music):
        """Return average distance between song_id and music library

        inputs: list of song_ids, a single song_id, and a music dictionary
        returns: average distance, computed as the sum of distances
                 divided by the number of distances considered, between song given
                 by song_id and the songs in song_id_list
        note: average_distance from empty list is 0
        >>> music = {'Stairway': [0,0],
        ...         '5th': [0,1],
        ...         'Blues': [1,1],
        ...         'Requiem': [1,0]}
        >>> average_distance([], 'Stairway', music)
        0.0
        >>> average_distance(['Stairway'], 'Stairway', music)
        0.0
        >>> average_distance(['5th'], 'Stairway', music)
        1.0
        >>> average_distance(['5th', 'Blues'], 'Stairway', music)
        1.5
        >>> average_distance(['5th', 'Blues', 'Requiem'], 'Stairway', music)
        1.3333333333333333
        """
        dist = 0.0
        for other in song_id_list:
            dist += distance(music[song_id], music[other])
        return dist/max(1, len(song_id_list))
```

```
In [7]: doctest.run_docstring_examples(average_distance, globals(), verbose=False)
```

Note that doctest using string comparisons between the expected and actual output, not more sophisticated tests like `==`. Thus in the above, we need 0.0 for expected return values, not 0.

Goodness Function

The "goodness" of a song is defined to be the average distance of the song from a list of disliked songs, minus the average distance of the song from a list of liked songs. This is meant to favor songs far away from disliked songs, but close to liked songs.

```
In [8]: def goodness(likes, dislikes, song_id, music):
        """Return `goodness` of song_id based on history of Likes/dislikes

        inputs: likes, dislikes are lists of 'liked' and 'disliked' song_ids.
                song_id is the id of a song we'd like to know the "goodness" of.
                music is a music dictionary.
        returns: "goodness" value (float) of song_id
        >>> music = {'Stairway': [0,0],
        ...         '5th': [0,1],
        ...         'Blues': [1,1],
        ...         'Requiem': [1,0]}
        >>> likes = []
        >>> dislikes = []
        >>> goodness(likes, dislikes, 'Stairway', music)
        0.0
        >>> likes = ['Requiem']
        >>> dislikes = ['5th', 'Blues']
        >>> goodness(likes, dislikes, 'Stairway', music)
        0.5
        >>> goodness(likes, dislikes, 'Back in Black', music)
        Traceback (most recent call last):
        ...
        KeyError: 'Back in Black'
        """
        return average_distance(dislikes, song_id, music) - \
               average_distance(likes, song_id, music)
```

```
In [9]: doctest.run_docstring_examples(goodness, globals(), verbose=False)
```

Next Song

Now to answer the key question -- what song should be picked next, based on previously played song likes and dislikes?

```

In [10]: def next_song(likes, dislikes, music):
        """Return next song to play based on history of likes/dislikes

        inputs: likes is list of 'liked' previously played song ids.
                 dislikes is list of 'disliked' previously played song ids.
                 music is a music dictionary.
        returns: ID for an unplayed song in dictionary with best goodness value

        >>> music = {'Stairway': [0,0],
        ...         '5th': [0,1],
        ...         'Blues': [1,1],
        ...         'Requiem': [1,0]}
        >>> likes = []
        >>> dislikes = ['Blues']
        >>> next_song(likes, dislikes, music)
        'Stairway'

        >>> likes = ['Blues']
        >>> dislikes = []
        >>> nxt = next_song(likes, dislikes, music)
        >>> nxt == '5th' or nxt == 'Requiem'
        True
        """
        played = set(likes) | set(dislikes)

        best_song_id = None
        best_goodness = 0

        # consider all songs in music
        for song_id in music:
            # disregard songs that have played already
            if song_id in played:
                continue

            # what is the goodness of the song we're considering?
            g = goodness(likes, dislikes, song_id, music)

            # if song is better than best_goodness, update best_goodness
            #if g > best_goodness:
            if g > best_goodness or best_song_id is None: #FIXED
                best_song_id = song_id
                best_goodness = g

        # an alternative loop to the above, removing the internal if song_id check:
        # for song_id in set(music.keys()) - played:

        # at this point, considered all unplayed songs, and must've seen the best
        return best_song_id

```

```

In [11]: doctest.run_docstring_examples(next_song, globals(), verbose=False)

```

```
In [12]: # Let's debug this case a little more...
music = {'Stairway': [0,0],
        '5th': [0,1],
        'Blues': [1,1],
        'Requiem': [1,0]}
likes = ['Blues']
dislikes = []
print('next_song (best goodness) is song id:',
      next_song(likes, dislikes, music))
for song_id in music:
    print('goodness for song id', song_id, "=",
          goodness(likes, dislikes, song_id, music))

next_song (best goodness) is song id: 5th
goodness for song id Stairway = -2.0
goodness for song id 5th = -1.0
goodness for song id Blues = 0.0
goodness for song id Requiem = -1.0
```

So, we only had one of the possible return values in our test case. Maybe we should be more thorough. Might the following work?

```
In [13]: likes = ['Blues']
dislikes = []
next_song(likes, dislikes, music)
```

Out[13]: '5th'

Not really. So how deal with that ambiguous return value, using doctest? Can update the doctest to craft a smarter check on return values:

```
In [ ]: """
...
>>> likes = ['Blues']
>>> dislikes = []
>>> nxt = next_song(likes, dislikes, music)
>>> nxt == '5th' or nxt == 'Requiem'
True
"""
```

Using Python random module to generate large test cases

(Note -- in 6.009 we're not allowing `import random`; but this gives you the idea)

How do you check your output is correct?

```
In [15]: import random
```

```
In [16]: def generate_test(music_size, gene_length, like_size, dislike_size):
    random.seed(6009) # A fixed random seed results in deterministic output.
    music = {}
    for i in range(music_size):
        music["song_"+str(i)] = [random.randint(0,1) for _ in range(gene_length)]
    likes = random.sample(music.keys(), like_size)
    dislikes = random.sample(music.keys(), dislike_size)
    return music, likes, dislikes
```

```
In [17]: music, likes, dislikes = generate_test(30, 10, 10, 0)
print("music:", music)
print("Example: song_1 =", music["song_1"])

next_song_id = next_song(likes, dislikes, music)
print("next_song_id:", next_song_id)
print("goodness:", goodness(likes, dislikes, next_song_id, music))
```

```
music: {'song_0': [0, 1, 0, 0, 1, 1, 0, 0, 1, 0], 'song_1': [1, 1, 0, 1, 1, 0, 1, 0, 1, 1], 'song_2': [1, 0, 1, 0, 1, 1, 0, 1, 0, 0], 'song_3': [0, 0, 0, 1, 1, 0, 1, 1, 0, 0], 'song_4': [0, 1, 0, 1, 1, 0, 1, 0, 1, 0], 'song_5': [1, 0, 1, 0, 0, 0, 1, 0, 1, 1], 'song_6': [0, 0, 0, 1, 0, 0, 0, 1, 1, 1], 'song_7': [1, 0, 0, 1, 0, 1, 0, 1, 1, 0], 'song_8': [0, 0, 0, 1, 1, 0, 1, 1, 1, 1], 'song_9': [1, 0, 1, 1, 1, 0, 0, 0, 0, 1], 'song_10': [0, 0, 0, 0, 1, 1, 1, 0, 1, 0], 'song_11': [0, 1, 1, 1, 0, 1, 1, 0, 0, 0], 'song_12': [0, 0, 1, 0, 0, 1, 1, 1, 1, 1], 'song_13': [1, 1, 0, 0, 0, 1, 1, 0, 0, 1], 'song_14': [0, 0, 0, 1, 0, 1, 1, 0, 0, 0], 'song_15': [0, 0, 0, 1, 0, 0, 0, 0, 0, 0], 'song_16': [0, 1, 0, 0, 1, 0, 1, 1, 0, 0], 'song_17': [0, 1, 1, 0, 1, 0, 0, 0, 0, 0], 'song_18': [0, 1, 1, 1, 1, 1, 0, 0, 1, 0], 'song_19': [1, 0, 1, 1, 1, 0, 0, 1, 0, 1], 'song_20': [0, 1, 0, 0, 0, 0, 1, 1, 1, 1], 'song_21': [1, 1, 0, 0, 1, 0, 0, 0, 0, 0], 'song_22': [0, 0, 1, 1, 0, 0, 1, 1, 0, 0], 'song_23': [1, 1, 1, 1, 0, 0, 0, 0, 0, 0], 'song_24': [1, 1, 0, 0, 0, 0, 0, 0, 0, 0], 'song_25': [0, 0, 0, 1, 0, 1, 0, 1, 1, 1], 'song_26': [1, 1, 0, 0, 1, 0, 0, 0, 1, 0], 'song_27': [1, 1, 1, 1, 1, 1, 1, 1, 0, 1], 'song_28': [1, 0, 1, 1, 1, 1, 1, 1, 1, 1], 'song_29': [0, 1, 1, 1, 0, 1, 1, 1, 1, 0]}
Example: song_1 = [1, 1, 0, 1, 1, 0, 1, 0, 1, 1]
next_song_id: song_4
goodness: -4.1
```

How do we know if the next_song_id is reasonable or correct?

```
In [18]: # Check what the best goodness actually is, among unplayed songs...
goodness_values = [goodness(likes, dislikes, song, music)
                    for song in music
                    if song not in likes and song not in dislikes]
max(goodness_values)

#or as a comprehension without building the list...
max(goodness(likes, dislikes, song, music)
    for song in music
    if song not in likes and song not in dislikes)
```

```
Out[18]: -4.1
```

Note that now we're **writing more complicated code to verify** or check answers. That's a job most likely better suited to unittest!